## Semaphores

## Andrews Chapter 04

- Chapter 3 used busy-waiting loops.
- Semaphores are easy to implement.
- Virtually all concurrent libraries include semaphores.
- Make critical sections easy.
- Can be used for signalling and scheduling.

```
sem mutex = 1;
process CS[i = 1 to n] {
  while (true) {
    P(mutex);
    critical section;
    V(mutex);
    noncritical section;
  }
}
```

**Figure 4.1** Semaphore solution to the critical section problem.

**Figure 4.2** Barrier synchronization using semaphores.

- A two-process barrier.
- Can use **count** to generalize (see Downey).
- Can use butterfly or dissemination-barrier.

```
typeT buf; /* a buffer of some type T */
sem empty = 1, full = 0;
process Producer[i = 1 to M] {
 while (true) {
    /* produce data, then deposit it in the buffer */
   P(empty);
   buf = data;
   V(full);
process Consumer[j = 1 to N] {
 while (true) {
   /* fetch result, then consume it */
   P(full);
   result = buf;
   V(empty);
```

Figure 4.3 Producers and consumers using semaphores.

• Split binary semaphores.

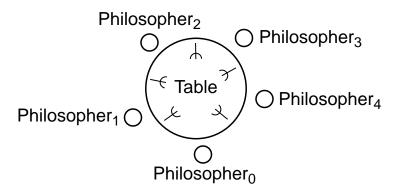
```
int front = 0, rear = 0;
sem empty = n, full = 0; /* n-2 \le mpty+full \le n */
process Producer {
 while (true) {
   produce message data and deposit it in the buffer;
   P(empty);
   buf[rear] = data; rear = (rear+1) % n;
   V(full);
process Consumer {
 while (true) {
   fetch message result and consume it;
   P(full);
   result = buf[front]; front = (front+1) % n;
   V(empty);
```

**Figure 4.4** Bounded buffer using semaphores.

• empty is initalized to n
Copyright © 2000 by Addison Wesley Longman, Inc.

```
int front = 0, rear = 0;
sem empty = n, full = 0; /* n-2 \le mpty+full \le n */
sem mutexD = 1, mutexF = 1; /* for mutual exclusion */
process Producer[i = 1 to M] {
 while (true) {
    . . .
   produce message data and deposit it in the buffer;
   P(empty);
   P(mutexD);
   buf[rear] = data; rear = (rear+1) % n;
   V(mutexD);
   V(full);
process Consumer[j = 1 to N] {
 while (true) {
   fetch message result and consume it;
   P(full);
   P(mutexF);
   result = buf[front]; front = (front+1) % n;
   V(mutexF);
   V(empty);
    • • •
```

**Figure 4.5** Multiple producers and consumers using semaphores.



**Figure 4.6** The dining philosophers.

```
process Philosopher[i = 0 to 4] {
   while (true) {
     think;
     acquire forks;
     eat;
     release forks;
   }
}
```

```
sem fork[5] = \{1, 1, 1, 1, 1\};
process Philosopher[i = 0 to 3] {
  while (true) {
    P(fork[i]); P(fork[i+1]); # get left fork then right
    eat;
    V(fork[i]); V(fork[i+1]);
    think;
process Philosopher[4] {
  while (true) {
    P(fork[0]); P(fork[4]); # get right fork then left
    eat;
    V(fork[0]); V(fork[4]);
    think;
```

**Figure 4.7** Dining philosophers solution using semaphores.

```
sem rw = 1;
process Reader[i = 1 to M] {
 while (true) {
    P(rw); # grab exclusive access lock
   read the database;
   V(rw); # release the lock
process Writer[j = 1 to N] {
 while (true) {
   P(rw); # grab exclusive access lock
   write the database;
   V(rw); # release the lock
```

**Figure 4.8** An overconstrained readers/writers solution.

```
int nr = 0;  # number of active readers
sem rw = 1;  # lock for reader/writer exclusion
process Reader[i = 1 to M] {
  while (true) {
    \langle nr = nr+1;
      if (nr == 1) P(rw); # if first, get lock
    read the database;
    \langle nr = nr-1;
      if (nr == 0) V(rw); # if last, release lock
process Writer[j = 1 to N] {
  while (true) {
    P(rw);
    write the database;
    V(rw);
```

**Figure 4.9** Outline of readers and writers solution.

```
int nr = 0;  # number of active readers
sem rw = 1;  # lock for access to the database
sem mutexR = 1; # lock for reader access to nr
process Reader[i = 1 to m] {
 while (true) {
    . . .
    P(mutexR);
      nr = nr+1;
      if (nr == 1) P(rw); # if first, get lock
   V(mutexR);
    read the database;
    P(mutexR);
      nr = nr-1;
      if (nr == 0) V(rw); # if last, release lock
   V(mutexR);
process Writer[j = 1 to n] {
 while (true) {
    . . .
    P(rw);
   write the database;
   V(rw);
```

**Figure 4.10** Readers and writers exclusion using semaphores.

```
int nr = 0, nw = 0;
## RW: (nr == 0 \vee nw == 0) \wedge nw <= 1
process Reader[i = 1 to m] {
  while (true) {
     \langle \text{await (nw == 0) nr = nr+1;} \rangle
     read the database;
     \langle nr = nr-1; \rangle
process Writer[j = 1 to n] {
  while (true) {
     \langle \text{await (nr == 0 and nw == 0) nw = nw+1;} \rangle
     write the database:
     \langle nw = nw-1; \rangle
```

Figure 4.11 A coarse-grained readers/writers solution.

- Waiting on *two* variables: no simple semaphore works.
- Our implementation of await with semaphores uses busy waiting.
- Is there a better way to simulate **await** with semaphores?

```
if (nw == 0 and dr > 0) {
   dr = dr-1; V(r); # awaken a reader, or
   }
elseif (nr == 0 and nw == 0 and dw > 0) {
   dw = dw-1; V(w); # awaken a writer, or
   }
else
   V(e); # release the entry lock
```

SIGNAL code for Figure 4.12.

- e is a mutex for the readers/writers.
- Instead of releasing **e**, signal *exactly one* of the semaphores.
- r, w, and e are a split binary semaphore: only one is nonzero.
- Passing the baton to one another.
- The signaller *also* decrements the counter.

```
## RW: (nr == 0 or nw == 0) and nw <= 1
int nr = 0,
    nw = 0;
sem e = 1, # controls entry to critical sections
    r = 0, # used to delay readers
    w = 0; # used to delay writers
               # at all times 0 \le (e+r+w) \le 1
int dr = 0, # number of delayed readers
              # number of delayed writers
    dw = 0;
process Reader[i = 1 to M] {
  while (true) {
    \# (await (nw == 0) nr = nr+1;)
      P(e);
      if (nw > 0) \{ dr = dr+1; V(e); P(r); \}
      nr = nr+1;
      SIGNAL:
                   # see text for details
    read the database;
    \# \langle nr = nr-1; \rangle
      P(e);
      nr = nr-1;
      SIGNAL;
process Writer[j = 1 to N] {
  while (true) {
    # \langle await (nr == 0 and nw == 0) nw = nw+1;\rangle
      if (nr > 0 \text{ or } nw > 0) \{ dw = dw+1; V(e); P(w); \}
      nw = nw+1;
      SIGNAL;
    write the database;
    \# \langle nw = nw-1; \rangle
      P(e);
      nw = nw-1;
      SIGNAL;
```

**Figure 4.12** Outline of readers and writers with passing the baton.

```
## RW: (nr==0 or nw==0) and nw<=1
int nr = 0,
    nw = 0;
sem e = 1, # controls entry to critical sections
    r = 0, # used to delay readers
    w = 0; # used to delay writers
               \# at all times 0 \le (e+r+w) \le 1
              # number of delayed readers
int dr = 0,
    dw = 0;
              # number of delayed writers
process Reader[i = 1 to M] {
  while (true) {
    \# (await (nw == 0) nr = nr+1;)
      P(e);
      if (nw > 0) \{ dr = dr+1; V(e); P(r); \}
      nr = nr+1;
      if (dr > 0) \{ dr = dr-1; V(r); \}
      else V(e);
    read the database;
    \# \langle nr = nr-1; \rangle
      P(e);
      nr = nr-1;
      if (nr == 0 \text{ and } dw > 0) \{ dw = dw-1; V(w); \}
      else V(e);
process Writer[j = 1 to N] {
  while (true) {
    # \langle \text{await (nr == 0 and nw == 0) nw = nw+1;} \rangle
      P(e);
      if (nr > 0 \text{ or } nw > 0) \{ dw = dw+1; V(e); P(w); \}
      nw = nw+1;
      V(e);
    write the database;
    \# \langle nw = nw-1; \rangle
      P(e);
      nw = nw-1;
      if (dr > 0) \{ dr = dr-1; V(r); \}
      elseif (dw > 0) { dw = dw-1; V(w); }
      else V(e);
                                      Some of the SIGNAL code can be simplified.
}
```

## Passing the baton more flexible

• Can give writers precedence by switching the **if** in **Writer**:

```
if (dw > 0) { dw = dw-1; V(w); }
elseif (dr > 0) { dr = dr - 1; V(r); }
else V(e);
```

• This change can be made without restructuring the code.

- Can force readers and writers to alternate when both are waiting:
  - delay a new reader when a writer is waiting
  - delay a new writer when a reader is waiting
  - awaken one waiting writer (if any) when a reader finishes
  - awaken all waiting readers (if any) when a writer finishes; otherwise awaken one waiting writer (if any)

```
bool free = true;
sem e = 1, b[n] = ([n] 0); # for entry and delay
typedef Pairs = set of (int, int);
Pairs pairs = \emptyset;
## SJN: pairs is an ordered set \land free \Rightarrow (pairs == \emptyset)
request(time,id):
  P(e);
  if (!free) {
    insert (time, id) in pairs;
   V(e); # release entry lock
    P(b[id]); # wait to be awakened
  free = false;
 V(e); # optimized since free is false here
release():
  P(e);
  free = true;
  if (P != \emptyset) {
    remove first pair (time, id) from pairs;
    V(b[id]); # pass baton to process id
  else V(e);
```

**Figure 4.14** Shortest-job-next allocation using semaphores.

```
#include <pthread.h>
                          /* standard lines */
#include <semaphore.h>
#define SHARED 1
#include <stdio.h>
void *Producer(void *); /* the two threads */
void *Consumer(void *);
                         /* global semaphores */
sem t empty, full;
                         /* shared buffer
int data;
                                               */
int numIters;
/* main() -- read command line and create threads */
int main(int argc, char *argv[]) {
 pthread t pid, cid;
                          /* thread and attributes */
 pthread attr t attr;
                                                    */
                          /* descriptors
 pthread attr init(&attr);
 pthread_attr_setscope(&attr, PTHREAD_SCOPE_SYSTEM);
  sem_init(&empty, SHARED, 1); /* sem empty = 1 */
  sem_init(&full, SHARED, 0); /* sem full = 0 */
 numIters = atoi(argv[1]);
 pthread create(&pid, &attr, Producer, NULL);
 pthread_create(&cid, &attr, Consumer, NULL);
 pthread join(pid, NULL);
  pthread_join(cid, NULL);
/* deposit 1, ..., numIters into the data buffer */
void *Producer(void *arg) {
  int produced;
  for (produced = 1; produced <= numIters; produced++) {</pre>
    sem_wait(&empty);
    data = produced;
    sem post(&full);
/* fetch numIters items from the buffer and sum them */
void *Consumer(void *arg) {
  int total = 0, consumed;
  for (consumed = 1; consumed <= numIters; consumed++) {</pre>
    sem wait(&full);
    total = total + data;
    sem post(&empty);
 printf("the total is %d\n", total);
```

Figure 4.15 Simple producer/consumer using Pthreads.