

Monitors

Andrews, Chapter 05

- More structure than semaphores.
- Can be implemented easily.
- Access to monitor variables is only through interface.
- Mutual exclusion of all monitor procedures is implicit:
procedures in the same monitor cannot be executed concurrently.
- Condition synchronization is by **condition variables**.

```

monitor FIFOsemaphore {
    int s = 0;  ## s >= 0
    cond pos;  #  signaled when s > 0

    procedure Psem() {
        if (s == 0)
            wait(pos);
        else
            s = s-1;
    }

    procedure Vsem() {
        if (empty(pos))
            s = s+1;
        else
            signal(pos);
    }
}

```

Figure 5.3 FIFO semaphore using passing the condition.

Copyright © 2000 by Addison Wesley Longman, Inc.

- Example monitor

Monitors

- Programs with monitors usually use two kinds of modules: active processes and passive monitors.
- Shared variables are inside the monitors.
- Processes communicate by calling procedures in the same monitor.
- Provided by Java.
- Provided in Unix.

Monitors

- The book uses a simple syntax for static monitors:

```
monitor mname {  
    declarations of permanent variables  
    initialization statements  
    procedures  
}
```

- And calling:

```
call mname.opname(arguments)
```

Three Properties of Monitors

- Only procedure names are visible outside the monitor.
- Monitors may not access variables declared outside the monitor.
- Permanent variables are initialized before any procedures are called.

Monitor Invariants

- Truth of the invariant should be established by the initialization.
- After any procedure is called, the invariant should remain true.

Monitor Mutual Exclusion and Synchronization

- Mutual exclusion is implicit:
no two procedures in a monitor execute concurrently.
- Synchronization is explicit:
uses **condition variables**.
- Condition variables are declared:

```
cond cv;
```

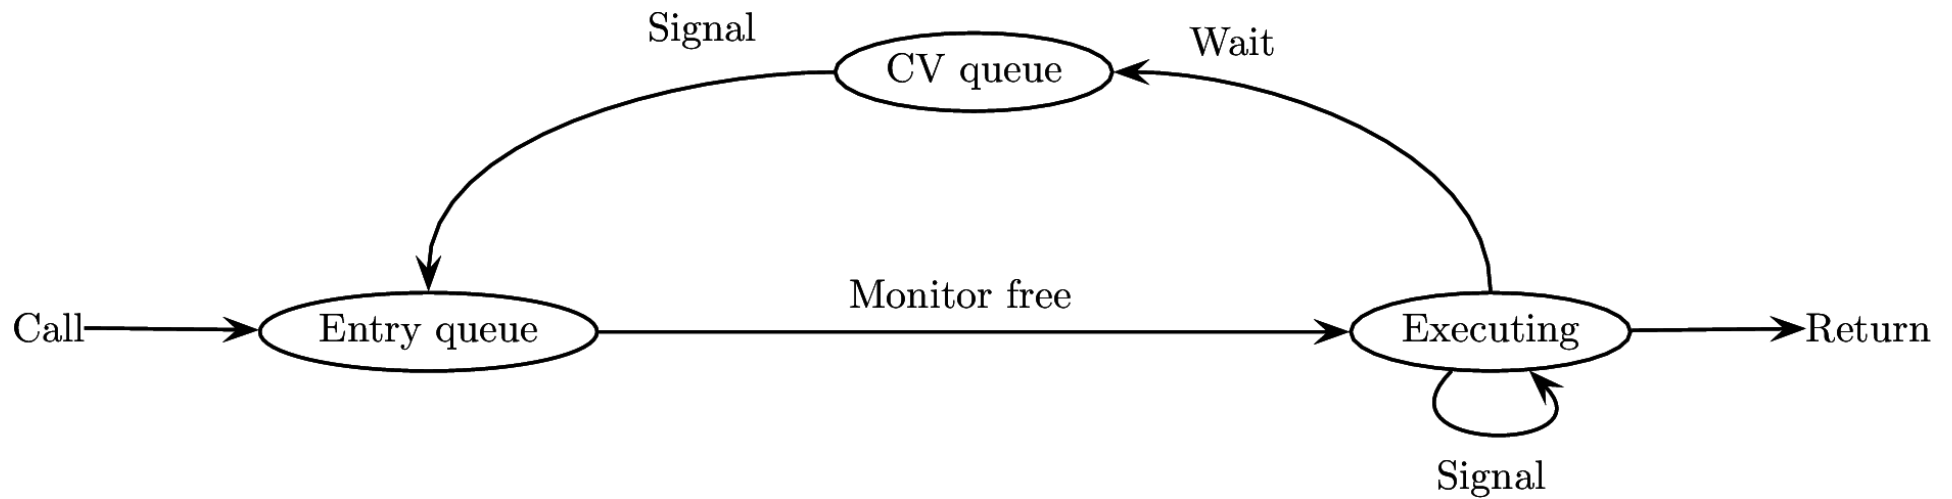
wait(cv)	wait at end of queue
wait(cv, rank)	wait in order of increasing value of rank
signal(cv)	awaken process at front of queue then continue
signal_all(cv)	awaken all processes on queue then continue
empty(cv)	true if wait queue is empty; false otherwise
minrank(cv)	value of rank of process at front of wait queue

Table 5.1 Operations on condition variables.

Copyright © 2000 by Addison Wesley Longman, Inc.

- Only called within the monitor.
- FIFO queue implicit, unless rank specified.


Signal and continue



Signal and continue

- The signaler continues and the signaled process waits.
- Nonpreemptive.

Signal and Wait



Signal and wait

- The signaler waits and the signaled process executes now.
- Preemptive.




Figure 5.1 State diagram for synchronization in monitors.

Copyright © 2000 by Addison Wesley Longman, Inc.

- Signal and continue used in Java and Pthreads.
- If signaled process put at front of entry queue:
signal and urgent wait

```

monitor Semaphore {
    int s = 0;  ## s >= 0
    cond pos;  # signaled when s > 0
    procedure Psem() {
        while (s == 0) wait(pos);
        s = s-1;
    }
    procedure Vsem() {
        s = s+1;
        signal(pos);
    }
}

```

Figure 5.2 Monitor implementation of a semaphore.

Copyright © 2000 by Addison Wesley Longman, Inc.

- Works for both SC and SW.
- SW is FIFO, but not SC.
- With SW, while can be replaced by if

```

monitor FIFOsemaphore {
    int s = 0;  ## s >= 0
    cond pos;   # signaled when s > 0

    procedure Psem() {
        if (s == 0)
            wait(pos);
        else
            s = s-1;
    }

    procedure Vsem() {
        if (empty(pos))
            s = s+1;
        else
            signal(pos);
    }
}

```

Figure 5.3 FIFO semaphore using passing the condition.

Differences between P and V and wait and signal

- wait and P:
 - wait *always* delays until a signal
 - P delays only if semaphore is zero
- signal and V:
 - signal has no effect on an empty queue
 - V either awakens a process or increments the semaphore

Book assumes Signal and Continue

- SC used in Unix, Java, and Pthreads.
- Makes `signal_all` well-defined.
- It is compatible with priority-based scheduling.
- It has simpler formal semantics.
- Historically, SW was first proposed for use in monitors.

```

monitor Bounded_Buffer {
    typeT buf[n];      # an array of some type T
    int front = 0,     # index of first full slot
        rear = 0;      # index of first empty slot
        count = 0;     # number of full slots
    ## rear == (front + count) % n
    cond not_full,     # signaled when count < n
        not_empty;    # signaled when count > 0

    procedure deposit(typeT data) {
        while (count == n) wait(not_full);
        buf[rear] = data; rear = (rear+1) % n; count++;
        signal(not_empty);
    }

    procedure fetch(typeT &result) {
        while (count == 0) wait(not_empty);
        result = buf[front]; front = (front+1) % n; count--;
        signal(not_full);
    }
}

```

Figure 5.4 Monitor implementation of a bounded buffer.


```

monitor RW_Controller {

    int nr = 0, nw = 0;  ## (nr == 0  $\vee$  nw == 0)  $\wedge$  nw <= 1
    cond oktoread;      # signaled when nw == 0
    cond oktowrite;     # signaled when nr == 0 and nw == 0

    procedure request_read() {
        while (nw > 0) wait(oktoread);
        nr = nr + 1;
    }

    procedure release_read() {
        nr = nr - 1;
        if (nr == 0) signal(oktowrite); # awaken one writer
    }

    procedure request_write() {
        while (nr > 0 || nw > 0) wait(oktowrite);
        nw = nw + 1;
    }

    procedure release_write() {
        nw = nw - 1;
        signal(oktowrite);      # awaken one writer and
        signal_all(oktoread);   # all readers
    }
}

```

Figure 5.5 Readers/writers solution using monitors.

```

monitor Shortest_Job_Next {
    bool free = true;  ## Invariant SJN: see text
    cond turn;          # signaled when resource available

    procedure request(int time) {
        if (free)
            free = false;
        else
            wait(turn, time);
    }

    procedure release() {
        if (empty(turn))
            free = true
        else
            signal(turn);
    }
}

```

Figure 5.6 Shortest-job-next allocation with monitors.

```

monitor Timer {
    int tod = 0;    ## invariant CLOCK -- see text
    cond check;    #   signaled when tod has increased

    procedure delay(int interval) {
        int wake_time;
        wake_time = tod + interval;
        while (wake_time > tod) wait(check);
    }

    procedure tick() {
        tod = tod + 1;
        signal_all(check);
    }
}

```

Figure 5.7 Interval timer with a covering condition.

```

monitor Timer {
    int tod = 0;    ## invariant CLOCK -- see text
    cond check;    #   signaled when minrank(check) <= tod
    procedure delay(int interval) {
        int wake_time;
        wake_time = tod + interval;
        if (wake_time > tod) wait(check, wake_time);
    }
    procedure tick() {
        tod = tod+1;
        while (!empty(check) && minrank(check) <= tod)
            signal(check);
    }
}

```

Figure 5.8 Interval timer with priority wait.

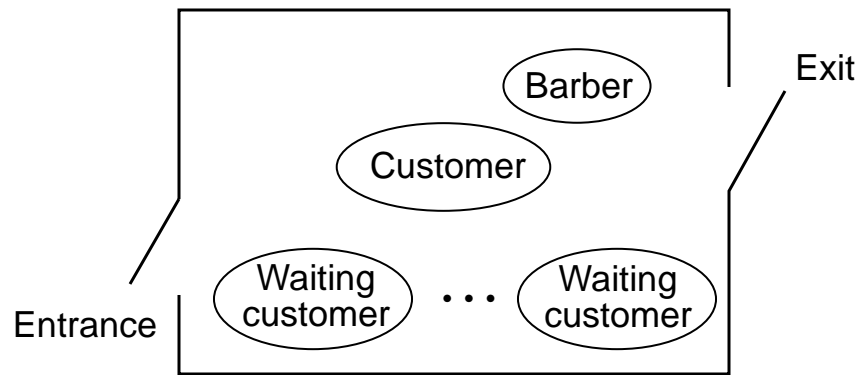


Figure 5.9 The sleeping barber problem.

Copyright © 2000 by Addison Wesley Longman, Inc.

```

monitor Barber_Shop {
    int barber = 0, chair = 0, open = 0;
    cond barber_available;      # signaled when barber > 0
    cond chair_occupied;        # signaled when chair > 0
    cond door_open;             # signaled when open > 0
    cond customer_left;         # signaled when open == 0

    procedure get_haircut() {
        while (barber == 0) wait(barber_available);
        barber = barber - 1;
        chair = chair + 1; signal(chair_occupied);
        while (open == 0) wait(door_open);
        open = open - 1; signal(customer_left);
    }

    procedure get_next_customer() {
        barber = barber + 1; signal(barber_available);
        while (chair == 0) wait(chair_occupied);
        chair = chair - 1;
    }

    procedure finished_cut() {
        open = open + 1; signal(door_open);
        while (open > 0) wait(customer_left);
    }
}

```

Figure 5.10 Sleeping barber monitor.

Skipping disk scheduler §5.3...except

Open and *Closed* monitors

- When a procedure in one monitor calls a procedure in a different monitor:
 - If it releases the original monitor it is an *open* call.
 - If it retains the lock on the original monitor it is a *closed* call.

Java: synchronized methods

```
class Counter {  
    private int c = 0;  
    public void increment() {  
        c++;  
    }  
    public int value() {  
        return c;  
    }  
}
```

```
class SynchronizedCounter {  
    private int c = 0;  
    public synchronized void increment() {  
        c++;  
    }  
    public synchronized int value() {  
        return c;  
    }  
}
```


Java: wait and notify

- Methods of the class Object
- Can only be used in synchronized methods
- Only a single delay queue per object
- There are no condition variables
 - exactly one (implied) condition variable per object
- FIFO not guaranteed
- notifyAll
- Nested calls are closed calls

```

// basic read or write; no exclusion
class RWbasic {
    protected int data = 0;  // the "database"
    public void read() {
        System.out.println("read:  " + data);
    }
    public void write() {
        data++;
        System.out.println("wrote:  " + data);
    }
}

class Reader extends Thread {
    int rounds;
    RWbasic RW;  // a reference to an RWbasic object
    public Reader(int rounds, RWbasic RW) {
        this.rounds = rounds;
        this.RW = RW;
    }
    public void run() {
        for (int i = 0; i < rounds; i++) {
            RW.read();
        }
    }
}

class Writer extends Thread {
    int rounds;
    RWbasic RW;
    public Writer(int rounds, RWbasic RW) {
        this.rounds = rounds;
        this.RW = RW;
    }
    public void run() {
        for (int i = 0; i < rounds; i++) {
            RW.write();
        }
    }
}

class Main {
    static RWbasic RW = new RWbasic();
    public static void main(String[] args) {
        int rounds = Integer.parseInt(args[0],10);
        new Reader(rounds, RW).start();
        new Writer(rounds, RW).start();
    }
}

```

```

// mutually exclusive read and write methods
class RWexclusive extends RWbasic {
    public synchronized void read() {
        System.out.println("read:  " + data);
    }
    public synchronized void write() {
        data++;
        System.out.println("wrote:  " + data);
    }
}

class Reader extends Thread {
    int rounds;
    RWexclusive RW;
    public Reader(int rounds, RWexclusive RW) {
        this.rounds = rounds;
        this.RW = RW;
    }
    public void run() {
        for (int i = 0; i < rounds; i++) {
            RW.read();
        }
    }
}

```

Exclusive readers/writers using Java.

```

// concurrent read or exclusive write
class ReadersWriters extends RWbasic {
    private int nr = 0;
    private synchronized void startRead() {
        nr++;
    }
    private synchronized void endRead() {
        nr--;
        if (nr == 0) notify(); // awaken waiting Writers
    }
    public void read() {
        startRead();
        System.out.println("read:  " + data);
        endRead();
    }
    public synchronized void write() {
        while (nr > 0)    // delay if any active Readers
            try { wait(); }
            catch (InterruptedException ex) {return;}
        data++;
        System.out.println("wrote:  " + data);
        notify();        // awaken another waiting Writer
    }
}

```

True readers/writers using Java.

Pthreads: locks and condition variables

Lock

```
...  
pthread_mutex_t mutex;  
...  
pthread_mutex_init(&mutex, NULL);  
...
```

Condition variable

```
...  
pthread_cond_t cv;  
...  
pthread_cond_init(&cv, NULL);  
...
```

Thread 0

```
...  
pthread_mutex_lock(&mutex);  
...  
pthread_cond_wait(&cv, &mutex);  
...  
pthread_mutex_unlock(&mutex);
```

Thread 1

```
...  
pthread_mutex_lock(&mutex);  
...  
pthread_cond_signal(&cv);  
...  
pthread_mutex_unlock(&mutex);
```

```

#include <pthread.h>
#include <stdio.h>
#define SHARED 1
#define MAXSIZE 2000 /* maximum matrix size */
#define MAXWORKERS 4 /* maximum number of workers */

pthread_mutex_t barrier; /* lock for the barrier */
pthread_cond_t go; /* condition variable */
int numWorkers; /* number of worker threads */
int numArrived = 0; /* number who have arrived */

/* a reusable counter barrier */
void Barrier() {
    pthread_mutex_lock(&barrier);
    numArrived++;
    if (numArrived < numWorkers)
        pthread_cond_wait(&go, &barrier);
    else {
        numArrived = 0; /* last worker awakens others */
        pthread_cond_broadcast(&go);
    }
    pthread_mutex_unlock(&barrier);
}

void *Worker(void *);
int size, stripSize; /* size == stripSize*numWorkers */
int sums[MAXWORKERS]; /* sums computed by each worker */
int matrix[MAXSIZE][MAXSIZE];

/* read command line, initialize, and create threads */
int main(int argc, char *argv[]) {
    int i, j;
    pthread_attr_t attr;
    pthread_t workerid[MAXWORKERS];

    /* set global thread attributes */
    pthread_attr_init(&attr);
    pthread_attr_setscope(&attr, PTHREAD_SCOPE_SYSTEM);

    /* initialize mutex and condition variable */
    pthread_mutex_init(&barrier, NULL);
    pthread_cond_init(&go, NULL);

    /* read command line */
    size = atoi(argv[1]);
    numWorkers = atoi(argv[2]);
    stripSize = size/numWorkers;

    /* initialize the matrix */
    for (i = 0; i < size; i++)
        for (j = 0; j < size; j++)
            matrix[i][j] = 1;

    /* create the workers, then exit main thread */
    for (i = 0; i < numWorkers; i++)
        pthread_create(&workerid[i], &attr,
            Worker, (void *) i);
}

```

```

    pthread_exit(NULL);
}

/* Each worker sums the values in one strip.
   After a barrier, worker(0) prints the total */
void *Worker(void *arg) {
    int myid = (int) arg;
    int total, i, j, first, last;

    /* determine first and last rows of my strip */
    first = myid*stripSize;
    last = first + stripSize - 1;

    /* sum values in my strip */
    total = 0;
    for (i = first; i <= last; i++)
        for (j = 0; j < size; j++)
            total += matrix[i][j];
    sums[myid] = total;
    Barrier();
    if (myid == 0) { /* worker 0 computes the total */
        total = 0;
        for (i = 0; i < numWorkers; i++)
            total += sums[i];
        printf("the total is %d\n", total);
    }
}

```

Figure 5.18 Parallel matrix summation using Pthreads.