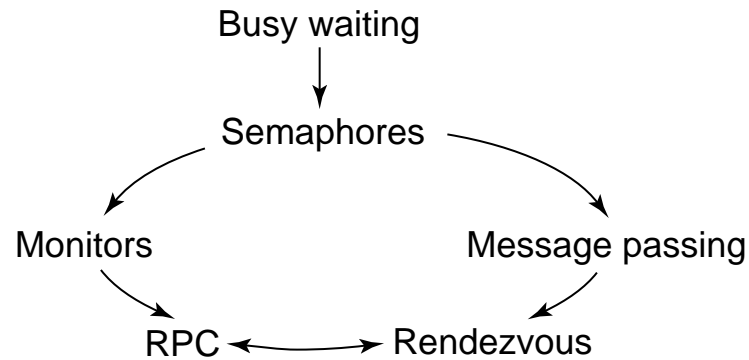


# Distributed Programming

- Distributed memory architectures.
- Concurrent programs here are usually called *distributed* programs.
- No shared variables.
  - No mutual exclusion necessary!
- Communicate and synchronize with *channels*:
  - one-way or two-way
  - synchronous (blocking) or asynchronous (nonblocking)
- Four basic mechanisms:
  - Chapter 7:
    - \* asynchronous message passing
    - \* synchronous message passing
  - Chapter 8:
    - \* RPC (remote procedure call)
    - \* rendezvous
- Chapter 9 describes several paradigms for distributed programming:
  - managers/workers, heartbeat, pipeline, probe/echo, broadcast, token passing, decentralized servers



Relationships between programming mechanisms.

Copyright © 2000 by Addison Wesley Longman, Inc.

- Monitors combine implicit exclusion with explicit signalling
- Message passing adds data to semaphore
- RPC and Rendezvous combine procedural interface of monitors with implicit message passing

# Message Passing

Andrews, Chapter 07

- Asynchronous message passing: channels are like semaphores.
- send and receive are like V and P
- receive is *blocking*
- May want to avoid blocking: `empty(ch)`
  - use with caution: may be unreliable
- The number of queued “messages” is the value of the semaphore.
- We assume messages are atomic and delivery is reliable and error-free.

```

chan input(char), output(char [MAXLINE]);
process Char_to_Line {
    char line[MAXLINE]; int i = 0;
    while (true) {
        receive input(line[i]);
        while (line[i] != CR and i < MAXLINE) {
            # line[0:i-1] contains the last i input characters
            i = i+1;
            receive input(line[i]);
        }
        line[i] = EOL;
        send output(line);
        i = 0;
    }
}

```

**Figure 7.1** Filter process to assemble lines of characters.

Copyright © 2000 by Addison Wesley Longman, Inc.

- Channels like this are called *mailboxes*

## Filters: Sorting

```
process Sort {  
    receive all numbers from channel input;  
    sort the numbers;  
    send the sorted numbers to channel output;  
}
```


- Suitable for “heavyweight” processes.
- Alternative: network of lightweight *merge* processes.

```

chan in1(int), in2(int), out(int);
process Merge {
    int v1, v2;
    receive in1(v1); # get first two input values
    receive in2(v2);
    # send smaller value to output channel and repeat
    while (v1 != EOS and v2 != EOS) {
        if (v1 <= v2)
            { send out(v1); receive in1(v1); }
        else # (v2 < v1)
            { send out(v2); receive in2(v2); }
    }
    # consume the rest of the non-empty input channel
    if (v1 == EOS)
        while (v2 != EOS)
            { send out(v2); receive in2(v2); }
    else # (v2 == EOS)
        while (v1 != EOS)
            { send out(v1); receive in1(v1); }
    # append a sentinel to the output channel
    send out(EOS);
}

```

**Figure 7.2** A filter process that merges two input streams.



**Figure 7.3** A sorting network of **Merge** processes.

Copyright © 2000 by Addison Wesley Longman, Inc.

- Original values can be lists sorted by “mediumweight” processes.

```

chan request(int clientID, types of input values);
chan reply[n](types of results);
process Server {
    int clientID;
    declarations of other permanent variables;
    initialization code;
    while (true) {    ## loop invariant MI
        receive request(clientID, input values);
        code from body of operation op;
        send reply[clientID](results);
    }
process Client[i = 0 to n-1] {
    send request(i, value arguments);    # "call" op
    receive reply[i](result arguments);    # wait for reply
}

```

**Figure 7.4** Clients and server with one operation.



```

type op_kind = enum(op1, ..., opn);
type arg_type = union(arg1, ..., argn);
type result_type = union(res1, ..., resn);
chan request(int clientID, op_kind, arg_type);
chan reply[n](res_type);

process Server {
    int clientID; op_kind kind; arg_type args;
    res_type results; declarations of other variables;
    initialization code;
    while (true) {      ## loop invariant MI
        receive request(clientID, kind, args);
        if (kind == op1)
            { body of op1; }
        ...
        else if (kind == opn)
            { body of opn; }
        send reply[clientID](results);
    }
}

process Client[i = 0 to n-1] {
    arg_type myargs; result_type myresults;
    place value arguments in myargs;
    send request(i, opj, myargs);      # "call" opj
    receive reply[i](myresults);        # wait for reply
}

```

**Figure 7.5** Clients and server with multiple operations.

```

monitor Resource_Allocator {
    int avail = MAXUNITS;
    set units = initial values;
    cond free;    # signaled when a process wants a unit
    procedure acquire(int &id) {
        if (avail == 0)
            wait(free);
        else
            avail = avail-1;
        remove(units, id);
    }
    procedure release(int id) {
        insert(units, id);
        if (empty(free))
            avail = avail+1;
        else
            signal(free);
    }
}

```

**Figure 7.6** Resource allocation monitor.

```

type op_kind = enum(ACQUIRE, RELEASE);
chan request(int clientID, op_kind kind, int unitid);
chan reply[n](int unitID);

process Allocator {
    int avail = MAXUNITS; set units = initial values;
    queue pending; # initially empty
    int clientID, unitID; op_kind kind;
    declarations of other local variables;
    while (true) {
        receive request(clientID, kind, unitID);
        if (kind == ACQUIRE) {
            if (avail > 0) { # honor request now
                avail--; remove(units, unitID);
                send reply[clientID](unitID);
            } else # remember request
                insert(pending, clientID);
        } else { # kind == RELEASE
            if empty(pending) { # return unitID to units
                avail++; insert(units, unitid);
            } else { # allocate unitID to a waiting client
                remove(pending, clientID);
                send reply[clientID](unitID);
            }
        }
    }
}

process Client[i = 0 to n-1] {
    int unitID;
    send request(i, ACQUIRE, 0) # "call" request
    receive reply[i](unitID);
    # use resource unitID, then release it
    send request(i, RELEASE, unitID);
    ...
}

```

**Figure 7.7** Resource allocator and clients.


*Monitor-Based Programs*

permanent variables  
procedure identifiers  
procedure call  
monitor entry  
procedure return  
**wait** statement  
**signal** statement  
procedure bodies

*Message-Based Programs*

local server variables  
**request** channel and operation kinds  
**send request(); receive reply**  
**receive request()**  
**send reply()**  
save pending request  
retrieve and process pending request  
arms of case statement on operation kind

**Table 7.1** Duality between monitors and message passing.



**Figure 7.8** Disk scheduling structures with message passing.

```

chan request(int clientID, int cyl, types of other arguments);
chan reply[n](types of results);

process Disk_Driver {
    queue left, right; # ordered queues of saved requests
    int clientID, cyl, headpos = 1, nsaved = 0;
    variables to hold other arguments in a request;
    while (true) {      ## loop invariant SST
        while (!empty(request) or nsaved == 0) {
            # wait for first request or receive another one
            receive request(clientID, cyl, ...);
            if (cyl <= headpos)
                insert(left, clientID, cyl, ...);
            else
                insert(right, clientID, cyl, ...);
            nsaved++;
        }
        # select best saved request from left or right
        if (size(left) == 0)
            remove(right, clientID, cyl, args);
        else if (size(right) == 0)
            remove(left, clientID, cyl, args);
        else
            remove request closest to headpos from left or right;
        headpos = cyl; nsaved--;
        access the disk;
        send reply[clientID](results);
    }
}

```

**Figure 7.9** Self-scheduling disk driver.

```

type kind = enum(READ, WRITE, CLOSE);
chan open(string fname; int clientID);
chan access[n](int kind, types of other arguments);
chan open_reply[m](int serverID); # server id or error
chan access_reply[m](types of results); # data, error, ...

process File_Server[i = 0 to n-1] {
    string fname; int clientID;
    kind k; variables for other arguments;
    bool more = false;
    variables for local buffer, cache, etc.;
    while (true) {
        receive open(fname, clientID);
        open file fname; if successful then:
        send open_reply[clientID](i); more = true;
        while (more) {
            receive access[i](k, other arguments);
            if (k == READ)
                process read request;
            else if (k == WRITE)
                process write request;
            else # k == CLOSE
                { close the file; more = false; }
            send access_reply[clientID](results);
        }
    }
}

process Client[j = 0 to m-1] {
    int serverID; declarations of other variables;
    send open("foo", j); # open file "foo"
    receive open_reply[j](serverID); # get back server id
    # use file then close it by executing the following
    send access[serverID](access arguments);
    receive access_reply[j](results);
    ...
}

```

**Figure 7.10** File servers and clients.

```

chan values(int), results[n](int smallest, int largest);

process P[0] {    # coordinator process
    int v;    # assume v has been initialized
    int new, smallest = v, largest = v; # initial state
    # gather values and save the smallest and largest
    for [i = 1 to n-1] {
        receive values(new);
        if (new < smallest)
            smallest = new;
        if (new > largest)
            largest = new;
    }
    # send the results to the other processes
    for [i = 1 to n-1]
        send results[i](smallest, largest)
}

process P[i = 1 to n-1] {
    int v;    # assume v has been initialized
    int smallest, largest;
    send values(v);
    receive results[i](smallest, largest);
}

```

**Figure 7.11** Exchanging values: centralized solution.



```

chan values[n](int);
process P[i = 0 to n-1] {
    int v;    # assume v has been initialized
    int new, smallest = v, largest = v; # initial state
    # send my value to the other processes
    for [j = 0 to n-1 st j != i]
        send values[j](v);
    # gather values and save the smallest and largest
    for [j = 1 to n-1] {
        receive values[i](new);
        if (new < smallest)
            smallest = new;
        if (new > largest)
            largest = new;
    }
}

```


**Figure 7.12** Exchanging values: symmetric solution.

```

chan values[n](int smallest, int largest);
process P[0] { # initiates the exchanges
    int v;    # assume v has been initialized
    int smallest = v, largest = v; # initial state
    # send v to next process, P[1]
    send values[1](smallest, largest);
    # get global smallest and largest from P[n-1] and
    #   pass them on to P[1]
    receive values[0](smallest, largest);
    send values[1](smallest, largest);
}
process P[i = 1 to n-1] {
    int v;    # assume v has been initialized
    int smallest, largest;
    # receive smallest and largest so far, then update
    #   them by comparing their values to v
    receive values[i](smallest, largest)
    if (v < smallest)
        smallest = v;
    if (v > largest)
        largest = v;
    # send the result to the next processes, then wait
    # to get the global result
    send values[(i+1) mod n](smallest, largest);
    receive values[i](smallest, largest);
}

```

**Figure 7.13** Exchanging values using a circular ring.



**Figure 7.14** Communication structures of the three programs.

Copyright © 2000 by Addison Wesley Longman, Inc.

```

channel values(int);

process Producer {
    int data[n];
    for [i = 0 to n-1] {
        do some computation;
        synch_send values(data[i]);
    }
}

process Consumer {
    int results[n];
    for [i = 0 to n-1] {
        receive values(results[i]);
        do some computation;
    }
}

```

Producer/consumer example using synchronous message passing.

Copyright © 2000 by Addison Wesley Longman, Inc.

```
channel in1(int), in2(int);  
process P1 {  
    int value1 = 1, value2;  
    synch_send in2(value1);  
    receive in1(value2);  
}  
  
process P2 {  
    int value1, value2 = 2;  
    synch_send in1(value2);  
    receive in2(value1);  
}
```

Exchanging values with synchronous message passing.

Copyright © 2000 by Addison Wesley Longman, Inc.

```

process GCD {
  int id, x, y;
  do true ->
    Client[*]?args(id, x, y); # input a "call"
    # repeat the following until x == y
    do x > y -> x = x - y;
    [] x < y -> y = y - x;
    od
    Client[id]!result(x); # return the result
  od
}

```

```

... GCD!args(i,v1,v2); GCD?result(r); ...

```

Greatest common divisor process and interface in CSP.

Copyright © 2000 by Addison Wesley Longman, Inc.

```

process Copy {    # one character buffer
    char c;
    do West?c -> East!c; od
}

```

```

process Copy {    # two character buffer
    char c1, c2;
    West?c1;
    do West?c2 -> East!c1; c1 = c2;
    [] East!c1 -> West?c1;
    od
}

```

```

process Copy {    # ten character buffer
    char buffer[10];
    int front = 0, rear = 0, count = 0;
    do count < 10; West?buffer[rear] ->
        count = count+1; rear = (rear+1) mod 10;
    [] count > 0; East!buffer[front] ->
        count = count-1; front = (front+1) mod 10;
    od
}

```

Versions of copy processes in CSP.

```

process Allocator {
    int avail = MAXUNITS;
    set units = initial values;
    int index, unitid;
    do avail > 0; Client[*]?acquire(index) ->
        avail--; remove(units, unitid);
        Client[index]!reply(unitid);
    [] Client[*]?release(index, unitid) ->
        avail++; insert(units,unitid);
    od
}

```

Resource allocator in CSP.

Copyright © 2000 by Addison Wesley Longman, Inc.



```
process P1 {  
    int value1 = 1, value2;  
    if P2!value1 -> P2?value2;  
    [] P2?value2 -> P2!value1;  
    fi  
}  
  
process P2 {  
    int value1, value2 = 2;  
    if P1!value2 -> P1?value1;  
    [] P1?value1 -> P1!value2;  
    fi  
}
```

Exchanging values in CSP.

Copyright © 2000 by Addison Wesley Longman, Inc.

```

process Sieve[1] {
    int p = 2;
    for [i = 3 to n by 2]
        Sieve[2]!i;    # pass odd numbers to Sieve[2]
}

process Sieve[i = 2 to L] {
    int p, next;
    Sieve[i-1]?p;      # p is a prime
    do Sieve[i-1]?next -> # receive next candidate
        if (next mod p) != 0 -> # if it might be prime,
            Sieve[i+1]!next;    # pass it on
        fi
    od
}

```

**Figure 7.15** Sieve of Eratosthenes in CSP.

```

CHAN OF BYTE comm :
PAR
    WHILE TRUE      -- keyboard input process
        BYTE ch :
        SEQ
            keyboard ? ch
            comm ! ch
    WHILE TRUE      -- screen output process
        BYTE ch :
        SEQ
            comm ? ch
            display ! ch

```

Producer/consumer example in Occam.

Copyright © 2000 by Addison Wesley Longman, Inc.

```

PROC Copy(CHAN OF BYTE West, Ask, East)
  BYTE c1, c2, dummy :
  SEQ
    West ? c1
    WHILE TRUE
      ALT
        West ? c2      -- West has a byte
        SEQ
          East ! c1
          c1 := c2
        Ask ? dummy    -- East wants a byte
        SEQ
          East ! c1
          West ? c1

```

Copy process in Occam.

Copyright © 2000 by Addison Wesley Longman, Inc.

```
COPY1 = West?c:char -> East!c -> COPY1
```

```
COPY = West?c1:char -> COPY2(c1)
```

```
COPY2(c1) =    West?c2:char -> East!c1 -> COPY2(c2)
              [ ]
              East!c1 -> West?c1:char -> COPY2(c1)
```

```
GCD = Input?id.x.y -> GCD(id, x, y)
```

```
GCD(id, x, y) = if (x = y) then
                  Output!id.x -> GCD
                else if (x > y) then
                  GCD(id, x-y, y)
                else
                  GCD(id, x, y-x)
```

Examples of Modern CSP.

Copyright © 2000 by Addison Wesley Longman, Inc.

```

#include "linda.h"
#define LIMIT 1000      /* upper bound for limit */

void worker() {
    int primes[LIMIT] = {2,3}; /* table of primes */
    int numPrimes = 1, i, candidate, isprime;

    /* repeatedly get candidates and check them */
    while(true) {
        if (RDP("stop")) /* check for termination */
            return;
        IN("candidate", ?candidate); /* get candidate */
        OUT("candidate", candidate+2); /* output next one */
        i = 0; isprime = 1;
        while (primes[i]*primes[i] <= candidate) {
            if (candidate%primes[i] == 0) { /* not prime */
                isprime = 0; break;
            }
            i++;
            if (i > numPrimes) { /* need another prime */
                numPrimes++;
                RD("prime", numPrimes, ?primes[numPrimes]);
            }
        }
        /* tell manager the result */
        OUT("result", candidate, isprime);
    }
}

real_main(int argc, char *argv[]) {
    int primes[LIMIT] = {2,3}; /* my table of primes */
    int limit, numWorkers, i, isprime;
    int numPrimes = 2, value = 5;
    limit = atoi(argv[1]); /* read command line */
    numWorkers = atoi(argv[2]);

    /* create workers and put first candidate in bag */
    for (i = 1; i <= numWorkers; i++)
        EVAL("worker", worker());
    OUT("candidate", value);

    /* get results from workers in increasing order */
    while (numPrimes < limit) {
        IN("result", value, ?isprime);
        if (isprime) { /* put value in table and TS */
            primes[numPrimes] = value;
            OUT("prime", numPrimes, value);
            numPrimes++;
        }
        value = value + 2;
    }
    /* tell workers to quit, then print the primes */
    OUT("stop");
    for (i = 0; i < limit; i++)
        printf("%d\n", primes[i]);
}

```

}

**Figure 7.16** Prime number generation in C-Linda.

Copyright © 2000 by Addison Wesley Longman, Inc.

```

#include <mpi.h>

main(int argc, char *argv[]) {
    int myid, otherid, size;
    int length = 1, tag = 1;
    int myvalue, othervalue;
    MPI_Status status;

    /* initialize MPI and get own id (rank) */
    MPI_Init(&argc, &argv);
    MPI_Comm_size(MPI_COMM_WORLD, &size);
    MPI_Comm_rank(MPI_COMM_WORLD, &myid);

    if (myid == 0) {
        otherid = 1; myvalue = 14;
    } else {
        otherid = 0; myvalue = 25;
    }

    MPI_Send(&myvalue, length, MPI_INT, otherid,
             tag, MPI_COMM_WORLD);
    MPI_Recv(&othervalue, length, MPI_INT, MPI_ANY_SOURCE,
             tag, MPI_COMM_WORLD, &status);
    printf("process %d received a %d\n", myid, othervalue);

    MPI_Finalize();
}

```

**Figure 7.17** MPI program to exchange values between two processes.



```

// Read a file and send it back to a client
import java.io.*; import java.net.*;

public class FileReaderServer {
    public static void main(String args[]) {
        try {
            // create server socket and
            // listen for connection on port 9999
            ServerSocket listen = new ServerSocket(9999);

            while (true) {
                System.out.println("waiting for connection");
                Socket socket = listen.accept(); // wait for client
                // create input and output streams to talk to client
                BufferedReader from_client =
                    new BufferedReader(new InputStreamReader
                        (socket.getInputStream()));
                PrintWriter to_client = new PrintWriter
                    (socket.getOutputStream());

                // get filename from client and check if it exists
                String filename = from_client.readLine();
                File inputFile = new File(filename);
                if (!inputFile.exists()) {
                    to_client.println("cannot open " + filename);
                    to_client.close(); from_client.close();
                    socket.close();
                    continue;
                }

                // read lines from filename and send to the client
                System.out.println("reading from file " + filename);
                BufferedReader input =
                    new BufferedReader(new FileReader(inputFile));
                String line;
                while ((line = input.readLine()) != null)
                    to_client.println(line);
                to_client.close(); from_client.close();
                socket.close();
            }
        } catch (Exception e) // report any exceptions
        { System.err.println(e); }
    }
}

```

**Figure 7.18** A file reader server in Java.

```

// Get file from RemoteFileServer and print on stdout
import java.io.*; import java.net.*;

public class Client {
    public static void main(String[] args) {
        try {
            // read command-line arguments
            if (args.length != 2) {
                System.out.println("need 2 arguments");
                System.exit(1);
            }
            String host = args[0];
            String filename = args[1];

            // open socket, then input and output streams to it
            Socket socket = new Socket(host,9999);
            BufferedReader from_server =
                new BufferedReader(new InputStreamReader
                    (socket.getInputStream()));
            PrintWriter to_server = new PrintWriter
                (socket.getOutputStream());

            // send filename to server, then read and print lines
            // until the server closes the connection
            to_server.println(filename); to_server.flush();
            String line;
            while ((line = from_server.readLine()) != null) {
                System.out.println(line);
            }
        }
        catch (Exception e)    // report any exceptions
        { System.err.println(e); }
    }
}

```

**Figure 7.19** A file reader client in Java.